



Machine Learning Classifier to predict sentiment for Amazon reviews using specific selected words.

author: bhavesh patel

```
In [1]: # Data set consist Amazon product reviews. We will use machine learning to udnerstand
# sentiment of each review. We will identify most positive and negative review for a given product.
# We will be using logistic regression as a classifier
# to predict the class of a discrete target variable (binary or multiclass) based on a model
# of class probability as a logistic function of a linear combination of the features.
# We use ROC curve (Receiver Operating Characteristic curve) for visulization.
# It is a plot of the true positive rate against the false positive rate for the different possible
# cutpoints of a diagnostic test.
```

```
In [2]: import graphlab
```

```
In [3]: # limit workers to preserve my laptop.
graphlab.set_runtime_config('GRAPHLAB_DEFAULT_NUM_PYLAMBDA_WORKERS', 4)
```

This non-commercial license of GraphLab Create for academic use is assigned to bhaveshhk8@gmail.com and will expire on October 17, 2017.

```
[INFO] graphlab.cython.cy_server: GraphLab Create v2.1 started. Logging: /tmp/graphlab_server_1479674229.log
```


```
In [4]: # now let's read amazon reviews.
product_reviews=graphlab.SFrame('amazon_baby.gl/')
```

```
In [5]: # lets browse the data.

# first show graphics locally here, not in a popup tab.
graphlab.canvas.set_target('ipynb')
```

```
In [6]: # data review using graph function.
```

```
product_reviews.show()
```

name		review		rating	
dtype:	str	dtype:	str	dtype:	float
num_unique (est.):	32,395	num_unique (est.):	185,979	num_unique (est.):	5
num_undefined:	284	num_undefined:	0	num_undefined:	0
frequent items:		frequent items:		min:	1
Vulli Sophie the ...		"		max:	5
Simple Wishes ...				median:	5
Infant Optics ...				mean:	4.12
Baby Einstein Take ...				std:	1.285
Cloud b Twilight ...				distribution of values:	
Fisher-Price ...					
Fisher-Price ...					
Graco Nautilus ...					
Leachco Snoogle ...					
Regalo Easy Step ...					
Baby Trend Diaper ...					
Skip Hop Zoo Pack ...					

```
In [7]: # remeber the defination of accuracy, which is defined as number of correct gueses over total data set records.  
# Let's add word count to the data set.
```

```
product_reviews['wordcount'] = graphlab.text_analytics.count_words(product_reviews['review'])
```

```
In [8]: # This time we are going to use specific worlds to create a model.  
selected_words = ['awesome', 'great', 'fantastic', 'amazing', 'love', 'wow',
```

```
        'horrible', 'bad', 'terrible', 'awful', 'hate']






# defining function which will allow to get count for a specific word in a dictionary.
# this will be used to create additional attributes in SFrame to use in classifier model.

def w_count(wdict, w):
    if w in wdict:
        count = wdict[w]
        return count
    else:
        return 0
```

```
In [9]: # Now let's create new attribute in product_reviews frame for each in the selected_words list.

for wd in selected_words:
    product_reviews[wd] = product_reviews['wordcount'].apply(lambda row: w_count(row,wd))
```

```
In [10]: product_reviews.show()
```







fantastic		amazing		love		wow		horrible		b
dtype:	int	dtype:	int	dtype:	int	dtype:	int	dtype:	int	d
num_unique (est.):	3	num_unique (est.):	5	num_unique (est.):	11	num_unique (est.):	5	num_unique (est.):	5	n
num_undefined:	0	num_undefined:	0	num_undefined:	0	num_undefined:	0	num_undefined:	0	n
min:	0	min:	0	min:	0	min:	0	min:	0	rr
max:	2	max:	4	max:	38	max:	6	max:	4	rr
median:	0	median:	0	median:	0	median:	0	median:	0	rr
mean:	0.005	mean:	0.007	mean:	0.229	mean:	7.846e-4	mean:	0.004	rr
std:	0.073	std:	0.09	std:	0.539	std:	0.032	std:	0.067	s
distribution of values:		distribution of values:		distribution of values:		distribution of values:		distribution of values:		di:
										

		8		
--	--	---	--	--

```
In [23]: # my model didn't work as it expected the target to be string or integer type.
# so I am converting rating to integer from float.

product_reviews['rating']=product_reviews['rating'].astype(int)

product_reviews.show()
```

fantastic		amazing		love		wow		horrible		b	
dtype:	int	dtype:	int	dtype:	int	dtype:	int	dtype:	int	dtype:	int
num_unique (est.):	3	num_unique (est.):	5	num_unique (est.):	11	num_unique (est.):	5	num_unique (est.):	5	num_unique (est.):	5
num_undefined:	0	num_undefined:	0	num_undefined:	0	num_undefined:	0	num_undefined:	0	num_undefined:	0
min:	0	min:	0	min:	0	min:	0	min:	0	min:	0
max:	2	max:	4	max:	38	max:	6	max:	4	max:	4
median:	0	median:	0	median:	0	median:	0	median:	0	median:	0
mean:	0.005	mean:	0.007	mean:	0.229	mean:	7.846e-4	mean:	0.004	mean:	0.004
std:	0.073	std:	0.09	std:	0.539	std:	0.032	std:	0.067	std:	0.067
distribution of values:		distribution of values:		distribution of values:		distribution of values:		distribution of values:		distribution of values:	
											

--	--	--	--	--	--

```
In [47]: # now we need to figure out sentiment. That is based on rating.
# There is column rating, which has 5 values. For now, we are going
# look into linear classifier which has binary value of 1 or 0.
# for that, we can define that any rating above 4 and 5 is positive aka 1
# any rating below 2 is negative aka 0.
# First, I don't like middle of the road rating 3, so ignore it.
product_reviews = product_reviews[product_reviews['rating'] != 3]

len(product_reviews)
```

Out[47]: 166752

```
In [50]: # now let's add directional column as we discussed above.

product_reviews['binrating'] = product_reviews['rating'] >= 4
```

```
In [51]: # let's create training and test dataset.

train_data, test_data = product_reviews.random_split(0.8, seed=0)
```

```
In [52]: # now let's create the classifier model.

selected_word_model = graphlab.logistic_classifier.create (train_data,
                                                            target='binrating',
                                                            features=selected_words,
                                                            validation_set=test_data)
```

Logistic regression:

Number of examples : 133448

Number of classes : 2

Number of feature columns : 11

Number of unpacked features : 11

Number of coefficients : 12

Starting Newton Method

+-----+-----+-----+-----+-----+				
Iteration	Passes	Elapsed Time	Training-accuracy	Validation-accuracy
+-----+-----+-----+-----+-----+				
1	2	0.210259	0.844299	0.842842
2	3	0.337249	0.844186	0.842842
3	4	0.461164	0.844276	0.843142
4	5	0.585145	0.844269	0.843142
5	6	0.705113	0.844269	0.843142
6	7	0.826983	0.844269	0.843142
+-----+-----+-----+-----+-----+				

SUCCESS: Optimal solution found.

```
In [53]: # now let's look at coefficients for each of the selected word.  
# sort them to understand which has most weightate in order.
```

```
coeff = selected_word_model['coefficients']  
coeff_sort=coeff.sort('value', ascending=False)  
  
print coeff_sort
```

+-----+-----+-----+-----+-----+					
name	index	class	value	stderr	
+-----+-----+-----+-----+-----+					
love	None	1	1.39989834302	0.0287147460124	
(intercept)	None	1	1.36728315229	0.00861805467824	
awesome	None	1	1.05800888878	0.110865296265	
amazing	None	1	0.892802422508	0.127989503231	
fantastic	None	1	0.891303090304	0.154532343591	
great	None	1	0.883937894898	0.0217379527921	

wow	None	1	-0.0541450123333	0.275616449416
bad	None	1	-0.985827369929	0.0433603009142
hate	None	1	-1.40916406276	0.0771983993506
awful	None	1	-1.76469955631	0.134679803365

[12 rows x 5 columns]

Note: Only the head of the SFrame is printed.

You can use `print_rows(num_rows=m, num_columns=n)` to print more rows and columns.

```
In [54]: # as we can see above, positive review with strong positive word make sense.
# let's look at most negative row.
```

```
coeff_sort[-1]
```

```
Out[54]: {'class': 1,
          'index': None,
          'name': 'terrible',
          'stderr': 0.09672419122285876,
          'value': -2.090499984872608}
```

```
In [56]: # Most negative make sense too. Let's look at all rows.
```

```
coeff_sort.print_rows(num_rows=12)
```

name	index	class	value	stderr
love	None	1	1.39989834302	0.0287147460124
(intercept)	None	1	1.36728315229	0.00861805467824
awesome	None	1	1.05800888878	0.110865296265
amazing	None	1	0.892802422508	0.127989503231
fantastic	None	1	0.891303090304	0.154532343591
great	None	1	0.883937894898	0.0217379527921
wow	None	1	-0.0541450123333	0.275616449416
bad	None	1	-0.985827369929	0.0433603009142
hate	None	1	-1.40916406276	0.0771983993506
awful	None	1	-1.76469955631	0.134679803365
horrible	None	1	-1.99651800559	0.0973584169028
terrible	None	1	-2.09049998487	0.0967241912229

[12 rows x 5 columns]

```
In [57]: # now let's evaluate the model to see how it performed.
```

```
selected_word_model.evaluate(test_data, metric='roc_curve')
```

```
Out[57]: {'roc_curve': Columns:
          threshold      float
          fpr          float
          tpr          float
          p            int
          n            int
```

Rows: 100001

Data:

threshold	fpr	tpr	p	n
0.0	1.0	1.0	27976	5328
1e-05	1.0	1.0	27976	5328
2e-05	1.0	1.0	27976	5328
3e-05	1.0	1.0	27976	5328
4e-05	1.0	1.0	27976	5328
5e-05	1.0	1.0	27976	5328
6e-05	1.0	1.0	27976	5328
7e-05	1.0	1.0	27976	5328
8e-05	1.0	1.0	27976	5328
9e-05	1.0	1.0	27976	5328


[100001 rows x 5 columns]

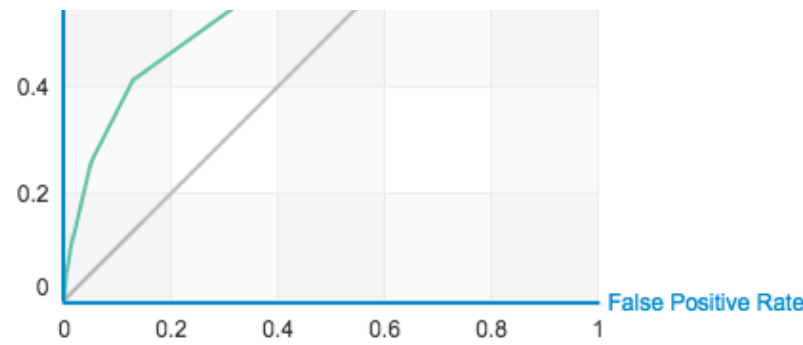
Note: Only the head of the SFrame is printed.

You can use `print_rows(num_rows=m, num_columns=n)` to print more rows and columns.

```
In [58]: selected_word_model.show(view='Evaluation')
```

Most recent model evaluation with dataset *test_data*





True Positive	False Negative	Accuracy	Precision
27836	140	0.843	0.845
False Positive	True Negative	Recall	F1 Score
5094	234	0.995	0.914

Threshold
0.501



AUC
0.665

now let's use this model to find out most positive/negative reviews for a product.

In [62]: *# get another SFrame for the product.*

```
diaper_champ_reviews = product_reviews[product_reviews['name']=='Baby Trend Diaper Champ']
```

In [63]: `len(diaper_champ_reviews)`

Out[63]: 298

In [69]: *# now that model is ready, let's use it.*


let's see how it predict each review sentiment for the diaper champ product.


we will add a column for each review. That will hold predicted sentiment by the model we built.


```
diaper_champ_reviews['predicted_sentiment_by_model']=selected_word_model.predict(diaper_champ_reviews, output_type='pro
```


```
diaper_champ_reviews.show()
```


great		fantastic		amazing		love		wow		h
dtype:	int	dtype:	int	dtype:	int	dtype:	int	dtype:	int	d
sum unique (set):	2	sum unique (set):	2	sum unique (set):	4	sum unique (set):	4	sum unique (set):	4	n


num_unique (est.):	3
num_undefined:	0
min:	0
max:	2
median:	0
mean:	0.198
std:	0.439
distribution of values:	
	

num_unique (est.):	2
num_undefined:	0
min:	0
max:	1
median:	0
mean:	0.003
std:	0.058
distribution of values:	
	

num_unique (est.):	1
num_undefined:	0
min:	0
max:	0
median:	0
mean:	0
std:	0
distribution of values:	
	

num_unique (est.):	4
num_undefined:	0
min:	0
max:	3
median:	0
mean:	0.312
std:	0.579
distribution of values:	
	

num_unique (est.):	1
num_undefined:	0
min:	0
max:	0
median:	0
mean:	0
std:	0
distribution of values:	
	

num_unique (est.):	1
num_undefined:	0
min:	0
max:	0
median:	0
mean:	0
std:	0
distribution of values:	
	

```
In [71]: # now let's figure out the most positive and negative review based on what the model predicted.
# let's sort it.
```

```
diaper_champ_reviews=diaper_champ_reviews.sort('predicted_sentiment_by_model', ascending=False)
```

```
In [73]: # most positive reviews:
diaper_champ_reviews[0]['review']
```

```
Out[73]: 'I LOVE LOVE LOVE this product! It is SO much easier to use than the Diaper Genie, (you need a PHD in poop to figure
out how to use the darn thing!) and it even takes the same bags as my kitchen trash can, which is super convenient,
and cost efficient as I can buy them in bulk.The only reason for not rating it a 5 star was that I did have one small
problem with it. The foam gasket in the barrell which keeps the poop smell inside the unit ripped somehow, and it
got VERY stinky. HOWEVER, I contacted the manufacturer through their website, and received an email back the same day
stating that this was unusual, and that replacement gaskets were on their way to me. They arrived inside of a week a
nd after replacing, it works great again! (They even sent me extras should it happen again)I HIGHLY recommend this di
```

aper pail over ANY competitors, you will not be sorry!'

In [74]: *# most negative review:*

```
diaper_champ_reviews[-1]['review']
```

Out[74]: "The Diaper Champ is TERRIBLE at keeping the smelly diapers from only smelling in the container. Our baby's room was constantly stinky (due to the Diaper Champ, not the baby!), and we were having to empty the container almost daily. What's the point of having a diaper disposal system if you can't dispose of diapers efficiently? Please don't buy this product unless you enjoy smelling those dirty diapers. The Diaper Champ just doesn't work."