



Predicting house value using machine learning regression analysis

author: bhavesh patel

credit: GraphLab

some definitions.

rmse = root mean squared error. This is used to identify errors and compare different models. rss = residual sum of squares is an error metric for regression.

These are two common measures of error regression, and RMSE is simply the square root of the mean RSS:

rmse = square root of (rss/n) where n=number of data points.

Predicting house value with regression analysis.

```
In [2]: import graphlab
```

```
In [3]: # Limit number of worker processes. This preserves system memory, which prevents hosted notebooks from crashing.
graphlab.set_runtime_config('GRAPHLAB_DEFAULT_NUM_PYLAMBDAS_WORKERS', 4)
```

```
[INFO] graphlab.cython.cy_server: GraphLab Create v2.1 started. Logging: /tmp/graphlab_server_1479486745.log
```

```
This non-commercial license of GraphLab Create for academic use is assigned to bhaveshhk8@gmail.com and will expire on October 17, 2017.
```

```
In [4]: # load data.
homesales = graphlab.SFrame('home_data.gl')
```

```
In [5]: homesales
```

Out[5]:

| | | | | | | | | |
|------------|---------------------------|---------|---|------|------|--------|---|---|
| 6414100192 | 2014-12-09 00:00:00+00:00 | 538000 | 3 | 2.25 | 2570 | 7242 | 2 | 0 |
| 5631500400 | 2015-02-25 00:00:00+00:00 | 180000 | 2 | 1 | 770 | 10000 | 1 | 0 |
| 2487200875 | 2014-12-09 00:00:00+00:00 | 604000 | 4 | 3 | 1960 | 5000 | 1 | 0 |
| 1954400510 | 2015-02-18 00:00:00+00:00 | 510000 | 3 | 2 | 1680 | 8080 | 1 | 0 |
| 7237550310 | 2014-05-12 00:00:00+00:00 | 1225000 | 4 | 4.5 | 5420 | 101930 | 1 | 0 |
| 1321400060 | 2014-06-27 00:00:00+00:00 | 257500 | 3 | 2.25 | 1715 | 6819 | 2 | 0 |
| 2008000270 | 2015-01-15 00:00:00+00:00 | 291850 | 3 | 1.5 | 1060 | 9711 | 1 | 0 |
| 2414600126 | 2015-04-15 00:00:00+00:00 | 229500 | 3 | 1 | 1780 | 7470 | 1 | 0 |
| 3793500160 | 2015-03-12 00:00:00+00:00 | 323000 | 3 | 2.5 | 1890 | 6560 | 2 | 0 |

| view | condition | grade | sqft_above | sqft_basement | yr_built | yr_renovated | zipcode | lat |
|------|-----------|-------|------------|---------------|----------|--------------|---------|-------------|
| 0 | 3 | 7 | 1180 | 0 | 1955 | 0 | 98178 | 47.51123398 |
| 0 | 3 | 7 | 2170 | 400 | 1951 | 1991 | 98125 | 47.72102274 |
| 0 | 3 | 6 | 770 | 0 | 1933 | 0 | 98028 | 47.73792661 |
| 0 | 5 | 7 | 1050 | 910 | 1965 | 0 | 98136 | 47.52082 |
| 0 | 3 | 8 | 1680 | 0 | 1987 | 0 | 98074 | 47.61681228 |
| 0 | 3 | 11 | 3890 | 1530 | 2001 | 0 | 98053 | 47.65611835 |
| 0 | 3 | 7 | 1715 | 0 | 1995 | 0 | 98003 | 47.30972002 |
| 0 | 3 | 7 | 1060 | 0 | 1963 | 0 | 98198 | 47.40949984 |
| 0 | 3 | 7 | 1050 | 730 | 1960 | 0 | 98146 | 47.51229381 |
| 0 | 3 | 7 | 1890 | 0 | 2003 | 0 | 98038 | 47.36840673 |

| long | sqft_living15 | sqft_lot15 |
|---------------|---------------|------------|
| -122.25677536 | 1340.0 | 5650.0 |
| -122.3188624 | 1690.0 | 7639.0 |
| -122.23319601 | 2720.0 | 8062.0 |

| | | |
|---------------|--------|----------|
| -122.39318505 | 1360.0 | 5000.0 |
| -122.04490059 | 1800.0 | 7503.0 |
| -122.00528655 | 4760.0 | 101930.0 |
| -122.32704857 | 2238.0 | 6819.0 |
| -122.31457273 | 1650.0 | 9711.0 |
| -122.33659507 | 1780.0 | 8113.0 |
| -122.0308176 | 2390.0 | 7570.0 |

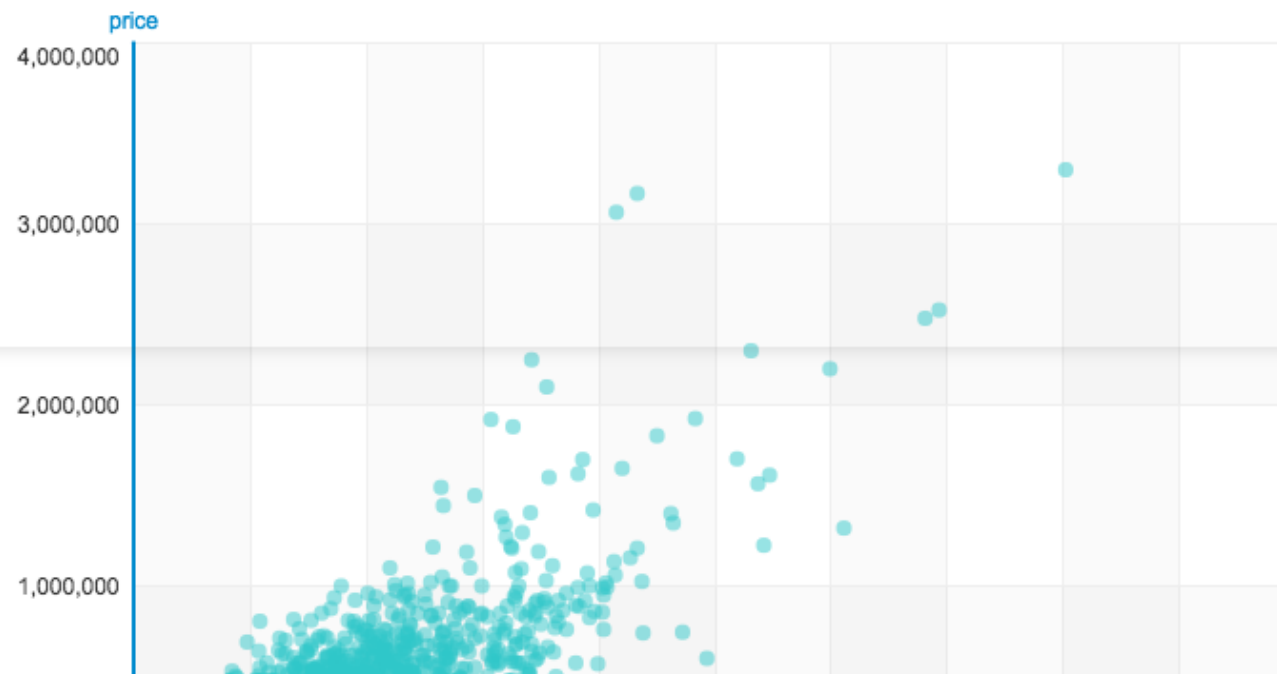
[21613 rows x 21 columns]

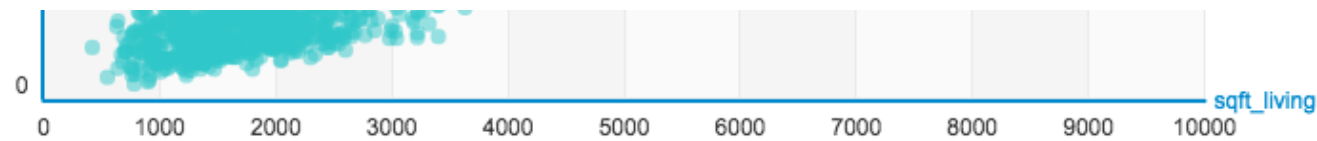
Note: Only the head of the SFrame is printed.

You can use `print_rows(num_rows=m, num_columns=n)` to print more rows and columns.

```
In [6]: # show graphis here, not a pop up window.
graphlab.canvas.set_target('ipynb')
```

```
In [7]: # there are total 21K+ rows. Let's view the data.
homesales.show(view="Scatter Plot", x="sqft_living", y="price")
```





Let's create regression model for prediction.

```
In [8]: # first get training data set and test data set. 80% training data, 20% test data.
train_data, test_data = homesales.random_split(0.8, seed=0)
```

```
In [9]: # build regression model with one variable for sq ft and store the results.
sqft_model = graphlab.linear_regression.create(train_data, target='price', features=['sqft_living'], validation_set=None)
```

Linear regression:

Number of examples : 17384

Number of features : 1

Number of unpacked features : 1

Number of coefficients : 2

Starting Newton Method

| Iteration | Passes | Elapsed Time | Training-max_error | Training-rmse |
|-----------|--------|--------------|--------------------|---------------|
| 1 | 2 | 1.014617 | 4349521.926170 | 262943.613754 |

```
In [10]: # let's perform some functions.
print test_data['price'].mean()
```

SUCCESS: Optimal solution found.

543054.042563

```
In [11]: print sqft_model.evaluate(test_data)
{'max_error': 4143550.8825285938, 'rmse': 255191.02870527358}
```

```
In [12]: print train_data['price'].mean()
539366.628221
```

```
In [13]: print sqft_model.evaluate(train_data)
{'max_error': 4349521.9261700595, 'rmse': 262943.6137536495}
```

```
In [14]: sqft_model.get('coefficients')
```

```
Out[14]:
```

| name | index | value | stderr |
|-------------|-------|----------------|---------------|
| (intercept) | None | -47114.0206702 | 4923.34437753 |
| sqft_living | None | 281.957850166 | 2.16405465323 |

[2 rows x 4 columns]

```
In [15]: sqft_model.show()
```

| Schema | |
|-----------------------------|-------|
| Number of coefficients | 2 |
| Number of examples | 17384 |
| Number of feature columns | 1 |
| Number of unpacked features | 1 |
| Hyperparameters | |
| L1 penalty | 0 |
| L2 penalty | 0.01 |
| Training Summary | |

| Training Summary | |
|-------------------------|----------------------------------|
| Solver | newton |
| Solver iterations | 1 |
| Solver status | SUCCESS: Optimal solution found. |
| Training time (sec) | 1.0201 |
| | |
| Settings | |
| Residual sum of squares | 1201918356336392 |
| Training RMSE | 262943.6138 |
| | |

Highest Positive Coefficients

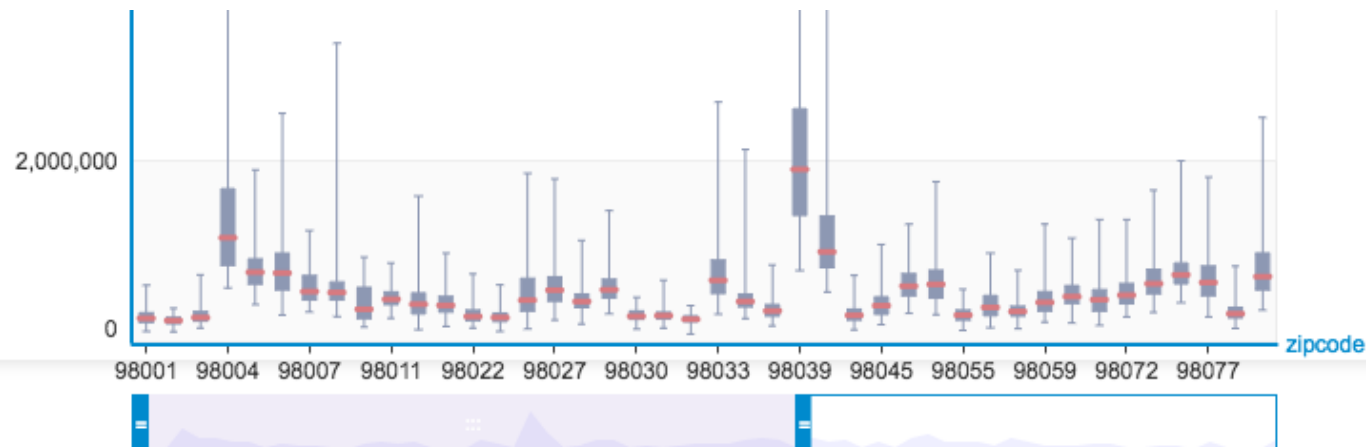
| | |
|-------------|---------|
| sqft_living | 281.958 |
|-------------|---------|

Lowest Negative Coefficients

| | |
|-------------|-------------|
| (intercept) | -47,114.021 |
|-------------|-------------|

```
In [16]: homesales.show(view='BoxWhisker Plot', x='zipcode', y='price')
```





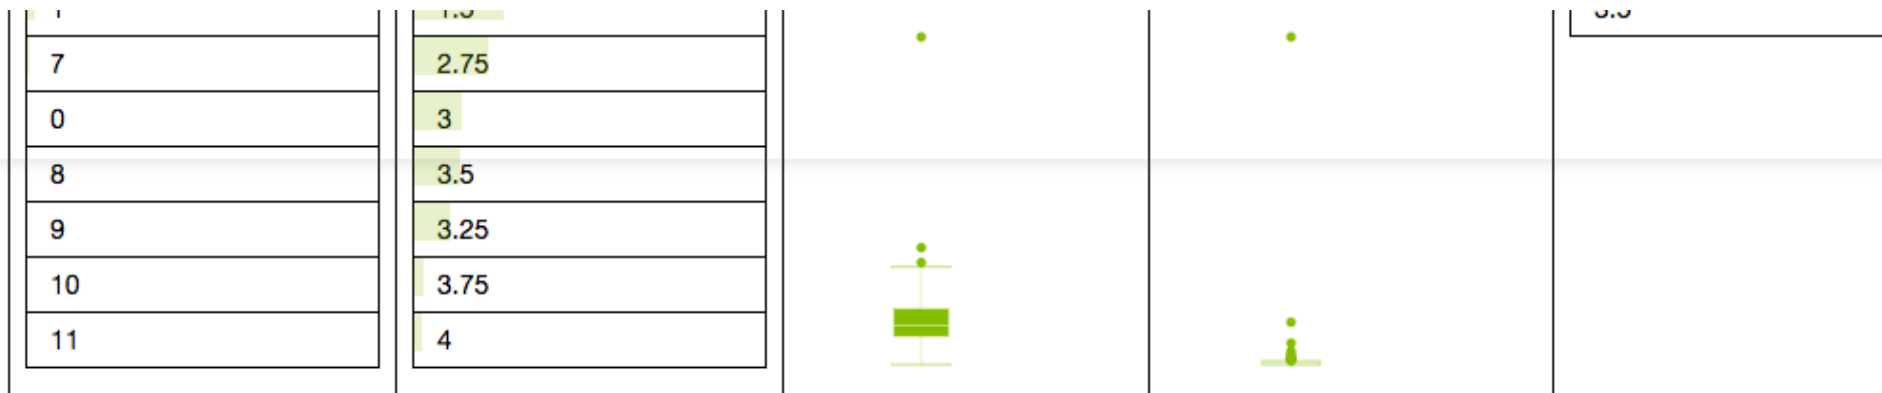
Explore more attributes in linear regression.

zip code 98039 has most expensive houses.

```
In [17]: my_features = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'zipcode']
```

```
In [18]: homesales[my_features].show()
```

| bedrooms | | bathrooms | | sqft_living | | sqft_lot | | floors | |
|--------------------|-----|--------------------|-----|-------------------------|-------|-------------------------|-------|--------------------|-----|
| dtype: | str | dtype: | str | dtype: | int | dtype: | int | dtype: | str |
| num_unique (est.): | 13 | num_unique (est.): | 30 | num_unique (est.): | 1,036 | num_unique (est.): | 9,747 | num_unique (est.): | 6 |
| num_undefined: | 0 | num_undefined: | 0 | num_undefined: | 0 | num_undefined: | 0 | num_undefined: | 0 |
| frequent items: | | frequent items: | | | | | | frequent items: | |
| 3 | | 2.5 | | | | | | 1 | |
| 4 | | 1 | | | | | | 2 | |
| 2 | | 1.75 | | | | | | 1.5 | |
| 5 | | 2.25 | | | | | | 3 | |
| 6 | | 2 | | | | | | 2.5 | |
| 1 | | 1.5 | | | | | | 3.5 | |
| | | | | distribution of values: | | distribution of values: | | | |



now let's build another regression model with more attributes.

```
In [19]: more_attribute_model=graphlab.linear_regression.create(train_data, target='price', features=my_features, validation_set=N
```

Linear regression:

Number of examples : 17384

Number of features : 6

Number of unpacked features : 6

Number of coefficients : 115

Starting Newton Method

```
+-----+-----+-----+-----+-----+
| Iteration | Passes | Elapsed Time | Training-max_error | Training-rmse |
+-----+-----+-----+-----+-----+
| 1         | 2      | 0.041900    | 3763208.270523    | 181908.848367 |
+-----+-----+-----+-----+-----+
```

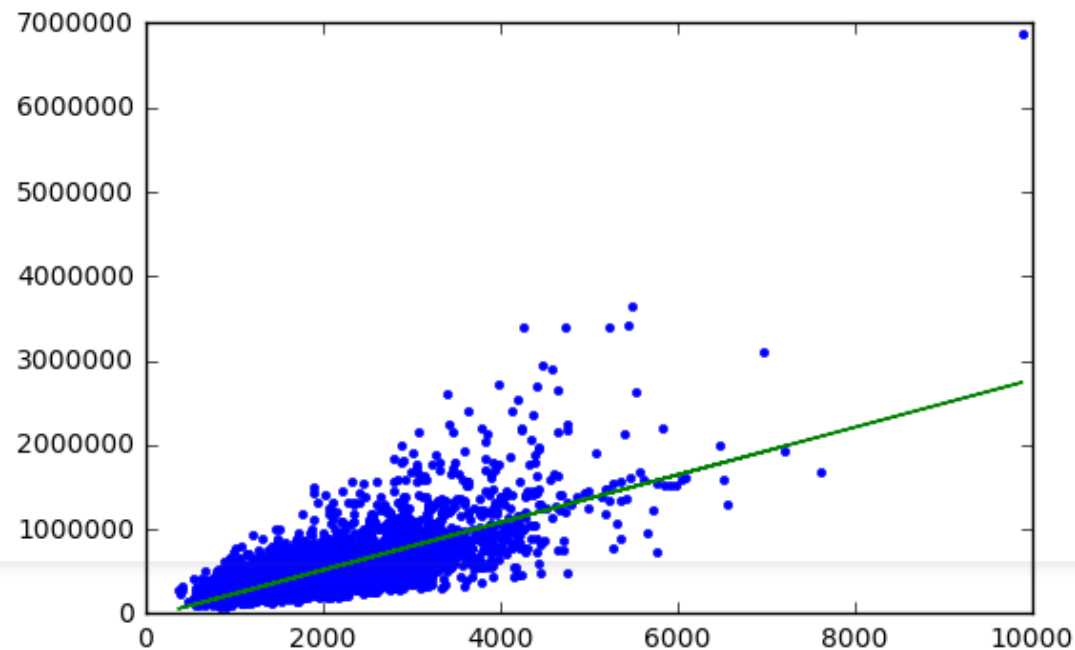
SUCCESS: Optimal solution found.


```
In [20]: # now let's compare the two model.  
print sqft_model.evaluate(test_data)  
print more_attribute_model.evaluate(test_data)  
  
{'max_error': 4143550.8825285938, 'rmse': 255191.02870527358}  
{'max_error': 3486584.509381705, 'rmse': 179542.4333126903}
```

```
In [21]: # with more attributes, we can see that error have reduced and RMSE has reduced too.  
# That's good and bad. We will see some overfitting in predicting house value.
```

```
In [22]: # let's get matplotlib to plot the regression analysis.  
import matplotlib.pyplot as plt  
%matplotlib inline  
  
plt.plot(test_data['sqft_living'],test_data['price'],'.',  
         test_data['sqft_living'],sqft_model.predict(test_data),'-')
```

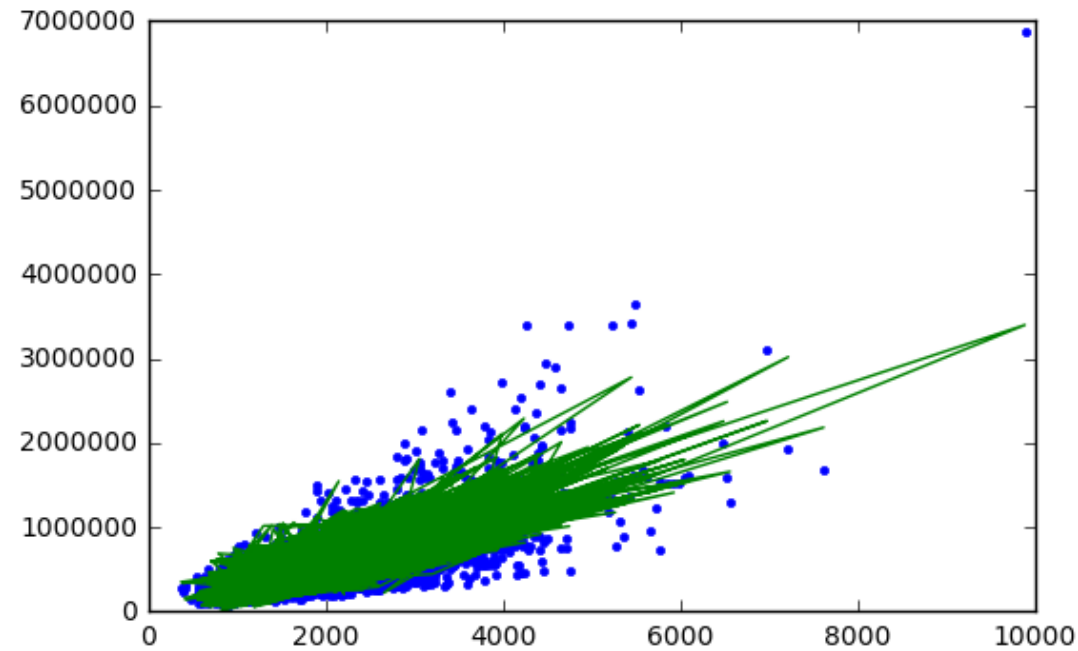
```
Out[22]: [<matplotlib.lines.Line2D at 0x11c822750>,  
          <matplotlib.lines.Line2D at 0x11c822810>]
```



```
In [23]: #Ok now let's build the model with second regression we built with more attributes.  
  
plt.plot(test_data['sqft_living'],test_data['price'],'.',
```

```
test_data['sqft_living'],more_attribute_model.predict(test_data),'-')
```

```
Out[23]: [<matplotlib.lines.Line2D at 0x11cb522d0>,  
<matplotlib.lines.Line2D at 0x11cb52390>]
```



```
In [30]: # now let's get one of the house id and try to predict the value of that house using both the regression models.  
# we need to get the object for one of the house. Let's pick the first house in the data, which has id  
# of 7129300520.
```

```
house1=homesales[homesales['id']=='7129300520']
```

```
In [31]: house1
```

```
Out[31]:
```

| id | date | price | bedrooms | bathrooms | sqft_living | sqft_lot | floors | waterfront |
|------------|---------------------------|--------|----------|-----------|-------------|----------|--------|------------|
| 7129300520 | 2014-10-13 00:00:00+00:00 | 221900 | 3 | 1 | 1180 | 5650 | 1 | 0 |

| view | condition | grade | sqft_above | sqft_basement | yr_built | yr_renovated | zipcode | lat |
|------|-----------|-------|------------|---------------|----------|--------------|---------|-------------|
| 0 | 3 | 7 | 1180 | 0 | 1955 | 0 | 98178 | 47.51123398 |

| long | sqft_living15 | sqft_lot15 |
|-------------|---------------|------------|
| 122.1450000 | 1180 | 5650 |

| | | |
|---------------|--------|--------|
| -122.25677536 | 1340.0 | 5650.0 |
|---------------|--------|--------|

[? rows x 21 columns]

Note: Only the head of the SFrame is printed. This SFrame is lazily evaluated.

You can use `sf.materialize()` to force materialization.

In [32]: *# now let's use this object and call our regression analysis model to predict house price.*

```
print sqft_model.predict(house1)
print more_attribute_model.predict(house1)
```

[285596.24252564885]

[211196.3344448047]

In []: *# ok. So actual house value is \$221,900. Our SQFT model predicted house value of \$285,596 while our more attribute model predicted house value of \$211,196. So we can say that more attribute value is showing us correct results. Well, its art with science. It is not always true. See below example for other way round.*

In [33]: *# now let's pick another house: 5309101200*

```
house2 = homesales[homesales['id']=='5309101200']
house2
```

Out[33]:

| id | date | price | bedrooms | bathrooms | sqft_living | sqft_lot | floors | waterfront |
|------------|---------------------------|--------|----------|-----------|-------------|----------|--------|------------|
| 5309101200 | 2014-06-05 00:00:00+00:00 | 620000 | 4 | 2.25 | 2400 | 5350 | 1.5 | 0 |

| view | condition | grade | sqft_above | sqft_basement | yr_built | yr_renovated | zipcode | lat |
|------|-----------|-------|------------|---------------|----------|--------------|---------|-------------|
| 0 | 4 | 7 | 1460 | 940 | 1929 | 0 | 98117 | 47.67632376 |

| long | sqft_living15 | sqft_lot15 |
|---------------|---------------|------------|
| -122.37010126 | 1250.0 | 4880.0 |

[? rows x 21 columns]

Note: Only the head of the SFrame is printed. This SFrame is lazily evaluated.

You can use `sf.materialize()` to force materialization.

In [34]: *# now let's predict values for this house.*

```
print sqft_model.predict(house2)
print more_attribute_model.predict(house2)
```

[629584.8197281545]

```
[721918.9333272863]
```

```
In [35]: # In this case, actual house value is: $620,000.  SQFT model predicted $629,584 while additional attribute model
# predicted $721,918.  Really bad!
```

```
In [38]: # now let's select the house with zip code

housesHighValueZip=homesales[homesales['zipcode']=='98039']
housesHighValueZip
```

```
Out[38]:
```

| | | | | | | | | |
|------------|---------------------------|---------|---|------|------|-------|---|---|
| 2540700110 | 2015-02-12 00:00:00+00:00 | 1905000 | 4 | 3.5 | 4210 | 18564 | 2 | 0 |
| 3262300940 | 2014-11-07 00:00:00+00:00 | 875000 | 3 | 1 | 1220 | 8119 | 1 | 0 |
| 3262300940 | 2015-02-10 00:00:00+00:00 | 940000 | 3 | 1 | 1220 | 8119 | 1 | 0 |
| 6447300265 | 2014-10-14 00:00:00+00:00 | 4000000 | 4 | 5.5 | 7080 | 16573 | 2 | 0 |
| 2470100110 | 2014-08-04 00:00:00+00:00 | 5570000 | 5 | 5.75 | 9200 | 35069 | 2 | 0 |
| 2210500019 | 2015-03-24 00:00:00+00:00 | 937500 | 3 | 1 | 1320 | 8500 | 1 | 0 |
| 6447300345 | 2015-04-06 00:00:00+00:00 | 1160000 | 4 | 3 | 2680 | 15438 | 2 | 0 |
| 6447300225 | 2014-11-06 00:00:00+00:00 | 1880000 | 3 | 2.75 | 2620 | 17919 | 1 | 0 |
| 2525049148 | 2014-10-07 00:00:00+00:00 | 3418800 | 5 | 5 | 5450 | 20412 | 2 | 0 |

| view | condition | grade | sqft_above | sqft_basement | yr_built | yr_renovated | zipcode | lat |
|------|-----------|-------|------------|---------------|----------|--------------|---------|-------------|
| 0 | 3 | 12 | 4860 | 0 | 1996 | 0 | 98039 | 47.61717049 |
| 0 | 3 | 11 | 4210 | 0 | 2001 | 0 | 98039 | 47.62060082 |
| 0 | 4 | 7 | 1220 | 0 | 1955 | 0 | 98039 | 47.63281908 |
| 0 | 4 | 7 | 1220 | 0 | 1955 | 0 | 98039 | 47.63281908 |
| 0 | 3 | 12 | 5760 | 1320 | 2008 | 0 | 98039 | 47.61512031 |
| 0 | 3 | 13 | 6200 | 3000 | 2001 | 0 | 98039 | 47.62888314 |
| 0 | 4 | 7 | 1320 | 0 | 1954 | 0 | 98039 | 47.61872888 |
| 2 | 3 | 8 | 2680 | 0 | 1902 | 1956 | 98039 | 47.61089438 |
| 1 | 4 | 9 | 2620 | 0 | 1949 | 0 | 98039 | 47.61435052 |
| 0 | 3 | 11 | 5450 | 0 | 2014 | 0 | 98039 | 47.62087993 |

| long | sqft_living15 | sqft_lot15 |
|---------------|---------------|------------|
| -122.23040939 | 3580.0 | 16054.0 |
| -122.2245047 | 3520.0 | 18564.0 |
| -122.23554392 | 1910.0 | 8119.0 |
| -122.23554392 | 1910.0 | 8119.0 |
| -122.22420058 | 3140.0 | 15996.0 |
| -122.23346379 | 3560.0 | 24345.0 |
| -122.22643371 | 2790.0 | 10800.0 |
| -122.22582388 | 4480.0 | 14406.0 |
| -122.22772057 | 3400.0 | 14400.0 |
| -122.23726918 | 3160.0 | 17825.0 |

[? rows x 21 columns]

Note: Only the head of the SFrame is printed. This SFrame is lazily evaluated.
You can use `sf.materialize()` to force materialization.

```
In [43]: # what's average price for the zip?

print housesHighValueZip['price'].mean()

2160606.6
```

```
In [47]: # SFrame filtering.
#select the houses that have 'sqft_living' higher than 2000 sqft but no larger than 4000 sqft.

houseFilter1 = homesales[(homesales['sqft_living'] > 2000) & (homesales['sqft_living'] <= 4000)]
houseFilter1
```

```
Out[47]:
```

| | | | | | | | | |
|------------|---------------------------|---------|---|------|------|-------|---|---|
| 1736800520 | 2015-04-03 00:00:00+00:00 | 662500 | 3 | 2.5 | 3560 | 9796 | 1 | 0 |
| 9297300055 | 2015-01-24 00:00:00+00:00 | 650000 | 4 | 3 | 2950 | 5000 | 2 | 0 |
| 2524049179 | 2014-08-26 00:00:00+00:00 | 2000000 | 3 | 2.75 | 3050 | 44867 | 1 | 0 |
| 7137970340 | 2014-07-03 00:00:00+00:00 | 285000 | 5 | 2.5 | 2270 | 6300 | 2 | 0 |

| | | | | | | | | |
|------------|---------------------------|--------|---|------|------|------|-----|---|
| 3814700200 | 2014-11-20 00:00:00+00:00 | 329000 | 3 | 2.25 | 2450 | 6500 | 2 | 0 |
| 1794500383 | 2014-06-26 00:00:00+00:00 | 937000 | 3 | 1.75 | 2450 | 2691 | 2 | 0 |
| 1873100390 | 2015-03-02 00:00:00+00:00 | 719000 | 4 | 2.5 | 2570 | 7173 | 2 | 0 |
| 8562750320 | 2014-11-10 00:00:00+00:00 | 580500 | 3 | 2.5 | 2320 | 3980 | 2 | 0 |
| 0461000390 | 2014-06-24 00:00:00+00:00 | 687500 | 4 | 1.75 | 2330 | 5000 | 1.5 | 0 |

| view | condition | grade | sqft_above | sqft_basement | yr_built | yr_renovated | zipcode | lat |
|------|-----------|-------|------------|---------------|----------|--------------|---------|-------------|
| 0 | 3 | 7 | 2170 | 400 | 1951 | 1991 | 98125 | 47.72102274 |
| 0 | 3 | 8 | 1860 | 1700 | 1965 | 0 | 98007 | 47.60065993 |
| 3 | 3 | 9 | 1980 | 970 | 1979 | 0 | 98126 | 47.57136955 |
| 4 | 3 | 9 | 2330 | 720 | 1968 | 0 | 98040 | 47.53164379 |
| 0 | 3 | 8 | 2270 | 0 | 1995 | 0 | 98092 | 47.32658071 |
| 0 | 4 | 8 | 2450 | 0 | 1985 | 0 | 98030 | 47.37386303 |
| 0 | 3 | 8 | 1750 | 700 | 1915 | 0 | 98119 | 47.63855772 |
| 0 | 3 | 8 | 2570 | 0 | 2005 | 0 | 98052 | 47.70732168 |
| 0 | 3 | 8 | 2320 | 0 | 2003 | 0 | 98027 | 47.5391103 |
| 0 | 4 | 7 | 1510 | 820 | 1929 | 0 | 98117 | 47.68228235 |

| long | sqft_living15 | sqft_lot15 |
|---------------|---------------|------------|
| -122.3188624 | 1690.0 | 7639.0 |
| -122.14529566 | 2210.0 | 8925.0 |
| -122.37541218 | 2140.0 | 4000.0 |
| -122.23345881 | 4110.0 | 20336.0 |
| -122.16892624 | 2240.0 | 7005.0 |
| -122.17228981 | 2200.0 | 6865.0 |
| -122.35985573 | 1760.0 | 3573.0 |
| -122.11029785 | 2630.0 | 6026.0 |
| -122.06971484 | 2580.0 | 3980.0 |


| | | |
|---------------|--------|--------|
| -122.36760203 | 1460.0 | 5000.0 |
|---------------|--------|--------|

[? rows x 21 columns]

Note: Only the head of the SFrame is printed. This SFrame is lazily evaluated.

You can use `sf.materialize()` to force materialization.

In [48]: `houseFilter1.show()`

| id | | date | | price | | bedrooms | | bathrooms | |
|--------------------|-------|-------------------------|----------|--|-------|--------------------|-----|--------------------|--|
| dtype: | str | dtype: | datetime | dtype: | int | dtype: | str | dtype: | |
| num_unique (est.): | 9,071 | num_unique (est.): | 353 | num_unique (est.): | 2,415 | num_unique (est.): | 12 | num_unique (est.): | |
| num_undefined: | 0 | num_undefined: | 0 | num_undefined: | 0 | num_undefined: | 0 | num_undefined: | |
| frequent items: | | frequent items: | | distribution of values: | | frequent items: | | frequent items: | |
| 5332200530 | | 2014-06-26 00:00:00+... | |  | | 4 | | 2.5 | |
| 9250900104 | | 2014-07-08 00:00:00+... | | | | 3 | | 2.25 | |
| 1254200015 | | 2014-07-29 00:00:00+... | | | | 5 | | 2.75 | |
| 1522059120 | | 2014-06-20 00:00:00+... | | | | 6 | | 1.75 | |
| 1561930020 | | 2015-04-21 00:00:00+... | | | | 2 | | 3 | |
| 1568100300 | | 2015-03-25 00:00:00+... | | | | 7 | | 2 | |
| 1630700361 | | 2014-07-01 00:00:00+... | | | | 8 | | 3.5 | |
| 1823059205 | | 2014-08-27 00:00:00+... | | | | 1 | | 3.25 | |
| 1825069031 | | 2014-06-23 00:00:00+... | | | | 9 | | 1.5 | |
| 1954420170 | | 2015-04-23 00:00:00+... | | | | 0 | | 1 | |
| 1974300020 | | 2015-04-27 00:00:00+... | | | | 10 | | 3.75 | |
| 2143700830 | | 2014-08-26 00:00:00+... | | | | 11 | | 4 | |

In [52]: `# alright, now to regression model 3, where we use lot more features.`

`even_more_features = [`

```

'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'zipcode',
'condition', # condition of house
'grade', # measure of quality of construction
'waterfront', # waterfront property
'view', # type of view
'sqft_above', # square feet above ground
'sqft_basement', # square feet in basement
'yr_built', # the year built
'yr_renovated', # the year renovated
'lat', 'long', # the lat-long of the parcel
'sqft_living15', # average sq.ft. of 15 nearest neighbors
'sqft_lot15', # average lot size of 15 nearest neighbors
]

```

```

In [53]: # build regression model using these features.
# more_attribute_model=graphlab.linear_regression.create(train_data, target='price', features=my_features, validation_se
even_more_att_model=graphlab.linear_regression.create(train_data, target='price', features=even_more_features, validation

```

Linear regression:

Number of examples : 17384

Number of features : 18

Number of unpacked features : 18

Number of coefficients : 127

Starting Newton Method

| | | | | |
|-----------|---------|--------------|--------------------|---------------|
| +-----+ | +-----+ | +-----+ | +-----+ | +-----+ |
| Iteration | Passes | Elapsed Time | Training-max_error | Training-rmse |
| +-----+ | +-----+ | +-----+ | +-----+ | +-----+ |
| 1 | 2 | 0.071567 | 3469012.450686 | 154580.940736 |
| +-----+ | +-----+ | +-----+ | +-----+ | +-----+ |

SUCCESS: Optimal solution found.


```
In [60]: # now let's get RMSE value for all three models we built.  
  
print sqft_model.evaluate(test_data)  
print more_attribute_model.evaluate(test_data)  
print even_more_att_model.evaluate(test_data)  
  
{'max_error': 4143550.8825285938, 'rmse': 255191.02870527358}  
{'max_error': 3486584.509381705, 'rmse': 179542.4333126903}  
{'max_error': 3556849.413858208, 'rmse': 156831.1168021901}
```